

PCRF Transformer Reliability Executive Report

Causal Reliability Flow & Derivative-Weighted Diagnostics Dashboard for QWEN/QWEN2.5-0.5B

1. Executive Summary & Core Gating Status

- **Target Architecture Profile:** QWEN/QWEN2.5-0.5B
- **Automated Path C Gating:** DISABLED (Discrete Exact Match accuracy utilized)
- **Composite PCRF Adoption Index:** 57.06 / 100
- **Strategic Deployment Directive:** RECOMMENDED WITH GATES: MONITORING & DATA CURATION ONLY (FALLBACK TUNING ACTIVE)
- **Production Security Risk Level:** ● [SAFE TO MEASURE & CURATE | DELAY INTERVENTION]

System Reliability Metrics At-a-Glance

Performance Parameter	Baseline (Pre-Intervention)	Post-Regularization	Net Gain / Loss Delta
Seen Validation Accuracy	45.00%	55.00%	+10.00%
Unseen Generalization Acc	45.00%	50.00%	+5.00%
Seen Validation Loss (NLL)	4.22905	3.68145	-0.54761
Unseen Validation Loss (NLL)	4.21506	3.86218	-0.35288
Unseen Perplexity (PPL)	67.70	47.57	-20.13

2. Integrated PCRF Scoreboard

Feature Track / Module	Baseline Value	PCRF Result Value	Track Score	Gating Status
Derivatives	0.00 (Unmeasured)	0.00364 (Avg Sensitivity)	3.6/100	SAFE_TO_APPLY
Curriculum Curation	Uniform Selection (Std=0.0)	PCRF Prioritized (Std=3.70)	74.0/100	SAFE_TO_APPLY
Structural Depth Monitor	Unmonitored Depth	Chain Reliability: 38.47%	38.5/100	DO_NOT_APPLY
Safe SFT Regularization	Unseen Acc: 45.0%	Unseen Acc: 50.0%	100.0/100	SAFE_TO_APPLY

3. Deep-Dive Diagnostics: Structural Bottlenecks & Representation Decay

Dynamic Decay Curve Insights for QWEN/QWEN2.5-0.5B

Causal analysis on QWEN/QWEN2.5-0.5B revealed the following trends across its structural segments:

- **Calculated System Chain Reliability (R_{sys}):** 38.4660%
- **Triggered Structural Flaw Heuristics:** None

Layer-wise Survival & Birnbaum Sensitivity Index

Layer	Survival Probability (r_l)	Birnbaum Derivative (D_R)
Block 00	88.85%	0.43291
Block 01	91.88%	0.41865
Block 02	99.25%	0.38755
Block 03	99.82%	0.38533
Block 04	99.84%	0.38527
Block 05	99.60%	0.38620
Block 06	99.62%	0.38613
Block 07	99.43%	0.38687
Block 08	99.36%	0.38712
Block 09	99.38%	0.38706
Block 10	99.52%	0.38650
Block 11	99.55%	0.38638
Block 12	99.55%	0.38638
Block 13	99.52%	0.38650
Block 14	99.33%	0.38724
Block 15	99.30%	0.38737
Block 16	99.21%	0.38774
Block 17	98.75%	0.38954
Block 18	98.53%	0.39041
Block 19	98.00%	0.39253
Block 20	97.15%	0.39592
Block 21	85.41%	0.45037
Block 22	83.94%	0.45827
Block 23	76.64%	0.50188

- **Dynamic Layer Integrity Trend Analysis:** U-Shaped Bottleneck Pattern (decay at boundaries)
 - Input Segment Average (First 15%): 93.33%
 - Highway Segment Average (Middle 70%): 99.19%
 - Output Segment Average (Last 15%): 82.00%

4. Empirical Perturbation Analysis: Layer-wise Sensitivity

- **Average System Sensitivity to Perturbation:** 0.00331
- **Highest Sensitivity Bottleneck Layer:** Layer 1 (Max Drift Delta: 0.01200)

Layer Index	Clean Target Prob	Perturbed Target Prob	Delta Prob
Layer 0	11.94%	10.89%	0.01049
Layer 1	11.94%	10.74%	0.01200
Layer 2	11.94%	11.49%	0.00457
Layer 3	11.94%	11.64%	0.00303
Layer 4	11.94%	11.53%	0.00417
Layer 5	11.94%	11.61%	0.00327
Layer 6	11.94%	12.69%	-0.00746
Layer 7	11.94%	11.96%	-0.00018
Layer 8	11.94%	12.00%	-0.00060
Layer 9	11.94%	11.51%	0.00435
Layer 10	11.94%	12.08%	-0.00140
Layer 11	11.94%	11.68%	0.00263
Layer 12	11.94%	12.49%	-0.00547
Layer 13	11.94%	11.92%	0.00021
Layer 14	11.94%	11.65%	0.00290
Layer 15	11.94%	11.96%	-0.00022
Layer 16	11.94%	11.55%	0.00394
Layer 17	11.94%	11.57%	0.00370
Layer 18	11.94%	11.95%	-0.00009
Layer 19	11.94%	11.58%	0.00367
Layer 20	11.94%	11.78%	0.00163
Layer 21	11.94%	11.78%	0.00166

Layer Index	Clean Target Prob	Perturbed Target Prob	Delta Prob
Layer 22	11.94%	11.97%	-0.00029
Layer 23	11.94%	11.80%	0.00144

5. Curriculum Prioritization & Semantic Error Concentration

The curriculum prioritization scoring system matches training prompts with a priority score:

$$\text{Priority Score} = \text{NLL} \times \left(1.0 + \sum_l \Delta_l \right)$$

This mathematical approach isolates low-impact prompts and prioritizes abstract syntax and logical structures. For example, QWEN/QWEN2.5-0.5B consistently struggles with advanced computer science and coding syntax prompts (high priority scores), while basic factual lookups (like capitals) remain stable.

Top 5 Highest Cascade-Risk Training Prompts

ID	Prompt	Target Completion	Priority Score
79	Complete with one word only: The reserved programming keyword used to initiate routine blocks in Python is	def	16.37
66	Complete with one word only: The digital counting framework representing information with 0 and 1 is	Binary	16.31
78	Complete with one word only: The reserved programming keyword used to declare structural blueprints in Python is	class	16.27
56	Complete with one word only: A liquid solution with a pH rating significantly higher than 7 is a	Base	15.64
74	Complete with one word only: The relational database directive used to fetch selected tuples from table arrays is	SELECT	15.34

Bottom 5 Lowest Cascade-Risk Training Prompts

ID	Prompt	Target Completion	Priority Score
40	Complete with one word only: The official capital city of Indonesia is	Jakarta	3.87
20	Complete with one word only: The official capital city of Thailand is	Bangkok	3.87
5	Complete with one word only: The official capital city of Japan is	Tokyo	3.63

ID	Prompt	Target Completion	Priority Score
29	Complete with one word only: The official capital city of Ireland is	Dublin	3.61
16	Complete with one word only: The official capital city of Argentina is	Buenos Aires	3.38

6. Root-Cause Debugging Hub: Failed Generations

The following table records the precise prompts where the model failed to generate an exact semantic match during validation evaluation. Use this trace to identify and rectify model logic regressions.

Split	ID	Prompt	Expected Target	Baseline Output	Candidate Output	Baseline NLL
seen_val	81	Complete with one word only: The official capital city of South Korea is	Seoul	`_____.		
A.`	Seoul. Seoul is the official	1.6797				
seen_val	84	Complete with one word only: The official capital city of Switzerland is	Bern	`_____.		
A.`	`_____.					
A.`	3.5469					
seen_val	86	Complete with one word only: The noble element designated by atomic number 10 is	Neon	`____		
A. Sodium`	`____					
A. Sodium`	9.9375					
seen_val	87	Complete with one word only: The	Sulfur	`____		

Split	ID	Prompt	Expected Target	Baseline Output	Candidate Output	Baseline NLL
		<i>volatile element designated by atomic number 16 is</i>				
A.		<i>Sodium` _____</i>				
A.	3.3125					
seen_val	89	<i>Complete with one word only: The yellow dwarf star supporting life at the center of our solar system is the</i>	Sun	closest star to Earth, and	closest star to Earth, and	3.2031
seen_val	90	<i>Complete with one word only: Mechanical acoustics are completely incapable of moving across a spatial</i>	Vacuum	medium. They are only able	medium. They are only able	12.5000
seen_val	96	<i>Complete with one word only: To enforce unique constraints with no duplicated items, algorithms utilize a</i>	Set	technique called "hashing."	technique called hashing. Hashing	6.5625
seen_val	97	<i>Complete with one word only: The hypermedia syntax used to format layout documents across the World Wide Web is</i>	HTML	`called _____.		
A`	`called _____.					
A`	3.8750					

Split	ID	Prompt	Expected Target	Baseline Output	Candidate Output	Baseline NLL
seen_val	98	Complete with one word only: An execution failure originating from incorrect program logic is called a	Bug	`(n) _____.		
A`	`(n) _____.					
A`	10.6875					
seen_val	99	Complete with one word only: A standardized text notation representing complex structural records is	JSON	`called a ____.		
A.`	`called a ____.					
A.`	10.4375					
...	...	...	...	...	...	(And 12 more trace details)

7. Actionable 4-Phase Enterprise Upgrade Strategy for QWEN/QWEN2.5-0.5B

To transform these automated insights into performance upgrades, the following engineering map is recommended for deployment:

Phase 1: Selective SFT Boundary Layer Regularization

Stop fine-tuning all model weights. Since intermediate layers are naturally self-correcting, freeze them. Concentrate training compute budgets exclusively on the high-risk boundary layers (Layers 0–1 and late projection blocks) to shrink training memory footprints by up to 50% while protecting historical parameter arrays from catastrophic forgetting.

Phase 2: Automated PCRF Dataset Curation

Deploy the curriculum_scores.csv output as a filtration gate before SFT training runs. Discard the bottom 30% to 50% of low-priority, high-redundancy training samples, saving massive amounts of GPU compute hours while ensuring the optimization loop is focused entirely on high-impact training samples.

Phase 3: Live Production "Canary" Monitors

Deploy the continuous Structural PCRF Plugin as an active endpoint wrapper. By feeding a tiny perturbed variant alongside the main user prompt, you can calculate the system chain reliability (R_{sys}) in real-time. If R_{sys} collapses below 75% for a given user query, route that query to a safe fallback model before a hallucinated or incorrect output reaches the user.

Phase 4: Automated CI/CD Governance Safety Gates

Integrate the Safe PCRF Controller as a hard gate in your Git-Ops pipelines. No custom-trained model should ever be promoted to your server registry unless it satisfies the continuous Path C non-inferiority constraints (protecting seen validation losses while ensuring a relative generalization gain of at least 5% on unseen splits).

8. Compute Environment & Host Profile Audit

- **Host Platform:** Linux 6.1.0-49-cloud-amd64
- **Detected CPU Cores:** 8
- **Host Memory Capacity:** 29.38 GB
- **GPU Platform Context:** Tesla T4 (14.56 GB VRAM)

Report programmatically generated by PCRF Reliability Suite v1.