PCRF Transformer Reliability Executive Report

Causal Reliability Flow & Derivative-Weighted Diagnostics Dashboard

---

1. Executive Summary & Core Gating Status

- **Target Architecture Profile:** `QWEN/QWEN2.5-0.5B`
- **Automated Path C Gating:** DISABLED (Discrete Exact Match accuracy utilized)
- **Composite PCRF Adoption Index:** `44.39 / 100`
- **Strategic Deployment Directive:** `REJECT ADOPTION: RESTORE BASELINE CONFIGURATIONS`
- **Production Security Risk Level:** 🔴 [HIGH OVER-STEERING RISK | BLOCK PARAMETER MUTATIONS]

System Reliability Metrics At-a-Glance

Performance Parameter	Baseline (Pre-Intervention)	Post-Regularization	Net Gain / Loss Delta
Seen Validation Accuracy	60.00%	65.00%	+5.00%
Unseen Generalization Acc	65.00%	60.00%	-5.00%
Seen Validation Loss (NLL)	4.54374	3.59827	-0.94547
Unseen Validation Loss (NLL)	4.26385	3.77604	-0.48781
Unseen Perplexity (PPL)	71.08	43.64	-27.44

---

2. Integrated PCRF Scoreboard

Feature Track / Module	Baseline Value	PCRF Result Value	Track Score	Gating Status
Derivatives	0.00 (Unmeasured)	0.00372 (Avg Sensitivity)	3.7/100	`SAFE_TO_APPLY`
Curriculum Curation	Uniform Selection (Std=0.0)	PCRF Prioritized (Std=3.78)	75.6/100	`SAFE_TO_APPLY`
Structural Depth Monitor	Unmonitored Depth	Chain Reliability: 35.11%	35.1/100	`SAFE_TO_APPLY`
Safe SFT Regularization	Unseen Acc: 65.0%	Unseen Acc: 60.0%	60.0/100	`MEASUREMENT_ONLY`

---

3. Deep-Dive Diagnostics: Structural Bottlenecks & Representation Decay

The "U-Shaped" Representation Decay Curve Insight

A major outcome of Scenario C is the discovery of the non-uniform nature of representation decay across the transformer's depth:

- **The Input Boundary (Layers 0-1):** Showcases immediate survival probability drops. Initial layers are vulnerable to input token embedding noise.
- **The Information Highway (Layers 2-20):** Shows exceptional stability with survival rates exceeding 90%. Representations here are robust and self-correcting.
- **The Output Boundary (Layers 21-23):** Drastically collapses (crashes to lowest levels), indicating that logit projection maps are highly chaotic and susceptible to minor drifts.
- **Calculated System Chain Reliability (R_{sys}):** 35.1096%
- **Triggered Structural Flaw Heuristics:** None

Layer-wise Survival & Birnbaum Sensitivity Index

Layer	Survival Probability (r_l)	Birnbaum Derivative (D_R)
Block 00	88.87%	0.39507
Block 01	91.44%	0.38395
Block 02	97.85%	0.35879
Block 03	99.17%	0.35402
Block 04	99.03%	0.35453
Block 05	99.02%	0.35459
Block 06	98.92%	0.35492
Block 07	98.73%	0.35560
Block 08	98.73%	0.35560
Block 09	98.81%	0.35532
Block 10	98.76%	0.35549
Block 11	98.84%	0.35521
Block 12	98.84%	0.35521
Block 13	98.83%	0.35526
Block 14	98.64%	0.35594
Block 15	98.64%	0.35594
Block 16	98.47%	0.35657
Block 17	98.21%	0.35748
Block 18	97.92%	0.35857

Layer	Survival Probability (r_l)	Birnbaum Derivative (D_R)
Block 19	97.42%	0.36040
Block 20	96.61%	0.36340
Block 21	86.46%	0.40606
Block 22	85.08%	0.41265
Block 23	78.33%	0.44822

---

4. Empirical Perturbation Analysis: Layer-wise Sensitivity

- **Average System Sensitivity to Perturbation:** `0.00400`
- **Highest Sensitivity Bottleneck Layer:** Layer `0` (Max Drift Delta: `0.01215`)

Layer Index	Clean Accuracy	Perturbed Accuracy	Delta (Sensitivity)
Layer 0	9.70%	8.48%	0.01215
Layer 1	9.70%	8.52%	0.01179
Layer 2	9.70%	9.89%	-0.00198
Layer 3	9.70%	8.66%	0.01034
Layer 4	9.70%	9.36%	0.00332
Layer 5	9.70%	9.47%	0.00223
Layer 6	9.70%	10.31%	-0.00618
Layer 7	9.70%	10.24%	-0.00549
Layer 8	9.70%	10.18%	-0.00483
Layer 9	9.70%	9.55%	0.00150
Layer 10	9.70%	9.21%	0.00484
Layer 11	9.70%	10.02%	-0.00325
Layer 12	9.70%	9.52%	0.00176
Layer 13	9.70%	10.06%	-0.00367
Layer 14	9.70%	10.04%	-0.00349
Layer 15	9.70%	9.29%	0.00409
Layer 16	9.70%	9.70%	-0.00004
Layer 17	9.70%	9.48%	0.00221
Layer 18	9.70%	9.39%	0.00302
Layer 19	9.70%	9.35%	0.00349

Layer Index	Clean Accuracy	Perturbed Accuracy	Delta (Sensitivity)
Layer 20	9.70%	9.56%	0.00134
Layer 21	9.70%	9.57%	0.00122
Layer 22	9.70%	9.36%	0.00340
Layer 23	9.70%	9.66%	0.00040

---

5. Curriculum Prioritization & Semantic Error Concentration

The curriculum prioritization scoring system matches training prompts with a priority score:

$$\text{Priority Score} = \text{NLL} \times \left(1.0 + \sum_{l} \Delta_l\right)$$

This mathematical approach filters out low-impact prompts and prioritizes abstract syntax and logical structures. For example, Qwen consistently struggles with Python code keywords (high priority scores), while factual lookups (like country capitals) remain highly stable (low priority scores).

Top 5 Highest Cascade-Risk Training Prompts

ID	Prompt	Target Completion	Priority Score
79	*The reserved programming keyword used to initiate routine blocks in Python is*	`def`	**18.42**
43	*The element with atomic number 6 is*	`Carbon`	**18.34**
78	*The reserved programming keyword used to declare structural blueprints in Python is*	`class`	**16.87**
66	*The digital counting framework representing information with 0 and 1 is*	`Binary`	**16.68**
63	*To store keyed associative records with rapid O(1) lookup, developers choose a*	`Map`	**14.99**

Bottom 5 Lowest Cascade-Risk Training Prompts

ID	Prompt	Target Completion	Priority Score
19	*The official capital city of Turkey is*	`Ankara`	**4.27**
40	*The official capital city of Indonesia is*	`Jakarta`	**4.23**
29	*The official capital city of Ireland is*	`Dublin`	**4.04**
5	*The official capital city of Japan is*	`Tokyo`	**3.93**
16	*The official capital city of Argentina is*	`Buenos Aires`	**3.56**

6. Root-Cause Debugging Hub: Failed Generations

The following table records the precise prompts where the model failed to generate an exact semantic match during validation evaluation. Use this trace to identify and rectify model logic regressions.

Split	ID	Prompt	Expected Target	Actual Generation	NLL Value
seen_val	86	*The noble element designated by atomic number 10 is*	`Neon`	`____	
A. Sodium`	10.9375				
seen_val	87	*The volatile element designated by atomic number 16 is*	`Sulfur`	`____	
A. Fluor`	3.4583				
seen_val	89	*The yellow dwarf star supporting life at the center of our solar system is the*	`Sun`	`most distant star in the Milky`	3.2344
seen_val	90	*Mechanical acoustics are completely incapable of moving across a spatial*	`Vacuum`	`domain. The only way to`	13.5625
seen_val	96	*To enforce unique constraints with no duplicated items, algorithms utilize a*	`Set`	`technique called "hashing."`	7.0625
seen_val	98	*An execution failure originating from incorrect program logic is called a*	`Bug`	`____ failure.	
A. Logical`	11.1875				
seen_val	99	*A standardized text notation representing complex structural records is*	`JSON`	`used in the field of structural`	9.5625
seen_val	100	*The active keyword used to bind external packages into Python script scopes is*	`import`	`____	
A. from`	7.3438				

Split	ID	Prompt	Expected Target	Actual Generation	NLL Value
unseen_val	106	*Mammalian red blood cells are chemically responsible for transporting vital*	`Oxygen`	`substances in the body. Which`	10.0000
unseen_val	111	*An unexpected, completely unpredictable event is idiomatic described as out of the*	`blue`	`ordinary. It is a situation`	1.7422
...	...	...	...	...	*(And 20 more recorded validation failures)*

---

7. Actionable 4-Phase Enterprise Upgrade Strategy

To transform these automated insights into performance upgrades, the following engineering map is recommended for deployment:

Phase 1: Selective SFT Boundary Layer Regularization

Stop fine-tuning all model weights. Since intermediate layers (Layers 2–20) are naturally self-correcting, freeze them. Concentrate training compute budgets exclusively on the high-risk boundary layers (Layers 0–1 and 21–23) to shrink training memory footprints by up to 50% while protecting historical parameter arrays from catastrophic forgetting.

Phase 2: Automated PCRF Dataset Curation

Deploy the `curriculum_scores.csv` output as a filtration gate before SFT training runs. Discard the bottom 30% to 50% of low-priority, high-redundancy training samples, saving massive amounts of GPU compute hours while ensuring the optimization loop is focused entirely on high-impact training samples.

Phase 3: Live Production "Canary" Monitors

Deploy the continuous Structural PCRF Plugin as an active endpoint wrapper. By feeding a tiny perturbed variant alongside the main user prompt, you can calculate the system chain reliability (R_{sys}) in real-time. If R_{sys} collapses below 75% for a given user query, route that query to a safe fallback model before a hallucinated or incorrect output reaches the user.

Phase 4: Automated CI/CD Governance Safety Gates

Integrate the Safe PCRF Controller as a hard gate in your Git-Ops pipelines. No custom-trained model should ever be promoted to your server registry unless it satisfies the continuous Path C non-inferiority constraints (protecting seen validation losses while ensuring a relative generalization gain of at least 5% on unseen splits).

---

8. Compute Environment & Host Profile Audit

- **Host Platform:** `Linux 6.1.0-49-cloud-amd64`
- **Detected CPU Cores:** `8`
- **Host Memory Capacity:** `29.38 GB`
- **GPU Platform Context:** `Tesla T4 (14.56 GB VRAM)`

Report programmatically generated by PCRF Reliability Suite v1.